

The Singular Value Decomposition (SVD): Detailed Analysis

AI Assistant

November 18, 2025

I. Intuition: Geometric Interpretation

The Singular Value Decomposition (SVD) is a factorization of a complex $m \times n$ matrix \mathbf{A} into three simple, interpretable linear transformations. Geometrically, the action of \mathbf{A} on the unit sphere in the input space (\mathbb{R}^n) results in an ellipse (or ellipsoid) in the output space (\mathbb{R}^m).

- The singular values in Σ are the **lengths of the semi-axes** of this resulting ellipse. They measure the **"gain"** or **"amplification"** of the matrix along specific directions.
- The columns of \mathbf{V} define the **input directions** (the principal axes of the original unit sphere) that get mapped to the axes of the ellipse.
- The columns of \mathbf{U} define the **output directions** (the axes of the ellipse) where the transformation ends up.

Sequence of Transformation

The transformation can be visualized as a sequence of three operations:

1. **Rotation/Reflection (\mathbf{V}^T)**: Orient the input space axes (\mathbf{v}_i) along the directions of maximum variance.
2. **Scaling (Σ)**: Stretches or shrinks these oriented axes by factors equal to the singular values (σ_i).
3. **Rotation/Reflection (\mathbf{U})**: Orient the final resulting ellipse in the output space. The columns of \mathbf{U} define the axes (\mathbf{u}_i) of this final ellipsoid.

II. The Mathematical Decomposition

The SVD of an $m \times n$ matrix \mathbf{A} is given by:

$$\mathbf{A} = \mathbf{U}\Sigma\mathbf{V}^T \quad (1)$$

A. Defining the Components

- U**: The $m \times m$ matrix of **Left Singular Vectors** (\mathbf{u}_i). It is an orthogonal matrix ($\mathbf{U}^T\mathbf{U} = \mathbf{I}_m$), meaning its column vectors form an orthonormal basis for \mathbb{R}^m .
- V**: The $n \times n$ matrix of **Right Singular Vectors** (\mathbf{v}_i). It is also an orthogonal matrix ($\mathbf{V}^T\mathbf{V} = \mathbf{I}_n$), and its column vectors form an orthonormal basis for \mathbb{R}^n .
- Σ** : The $m \times n$ rectangular diagonal matrix of **Singular Values** (σ_i).

B. Derivation of Singular Values

The singular values are unique and always non-negative, ordered such that $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_r > 0$, where $r = \text{rank}(\mathbf{A})$. They are directly related to the eigenvalues of the square symmetric matrices $\mathbf{A}^T \mathbf{A}$ and $\mathbf{A} \mathbf{A}^T$.

- The singular values σ_i are the square roots of the non-zero eigenvalues (λ_i) of $\mathbf{A}^T \mathbf{A}$: $\sigma_i = \sqrt{\lambda_i(\mathbf{A}^T \mathbf{A})}$.
- The right singular vectors (\mathbf{v}_i) are the corresponding eigenvectors of $\mathbf{A}^T \mathbf{A}$.
- The left singular vectors (\mathbf{u}_i) are the corresponding eigenvectors of $\mathbf{A} \mathbf{A}^T$.

III. Key Insights: The Four Fundamental Subspaces

The rank r defines the size of the non-trivial subspaces. The columns of \mathbf{U} and \mathbf{V} naturally span the four fundamental subspaces:

Subspace	Basis (from SVD)	Dimension
Column Space $\text{Col}(\mathbf{A}) \subset \mathbb{R}^m$	$\{\mathbf{u}_1, \dots, \mathbf{u}_r\}$	r
Null Space $\text{Null}(\mathbf{A}) \subset \mathbb{R}^n$	$\{\mathbf{v}_{r+1}, \dots, \mathbf{v}_n\}$	$n - r$
Row Space $\text{Col}(\mathbf{A}^T) \subset \mathbb{R}^n$	$\{\mathbf{v}_1, \dots, \mathbf{v}_r\}$	r
Left Null Space $\text{Null}(\mathbf{A}^T) \subset \mathbb{R}^m$	$\{\mathbf{u}_{r+1}, \dots, \mathbf{u}_m\}$	$m - r$

IV. Low-Rank Approximation and Data Compression

The full expansion of \mathbf{A} is the sum of r rank-one matrices:

$$\mathbf{A} = \sigma_1 \mathbf{u}_1 \mathbf{v}_1^T + \sigma_2 \mathbf{u}_2 \mathbf{v}_2^T + \dots + \sigma_r \mathbf{u}_r \mathbf{v}_r^T = \sum_{i=1}^r \sigma_i \mathbf{u}_i \mathbf{v}_i^T \quad (2)$$

The power of SVD lies in the **Eckart–Young Theorem**, which states that the best possible rank k approximation (\mathbf{A}_k) is achieved by keeping only the k largest singular values and their corresponding vectors:

$$\mathbf{A}_k = \sum_{i=1}^k \sigma_i \mathbf{u}_i \mathbf{v}_i^T \quad \text{for } k < r \quad (3)$$

This property is critical for data analysis, image compression, and matrix denoising because the information is ordered by importance (magnitude of σ_i).